

# Here there be dragons

## The HPC Code Development Journey

Phil Tooley  
University of Sheffield

8 April 2019

scheduler memory  
libraries MPI user interface  
testing language  
performance CUDA bugs  
OpenMP documentation licensing  
throughput compilers  
version control algorithms

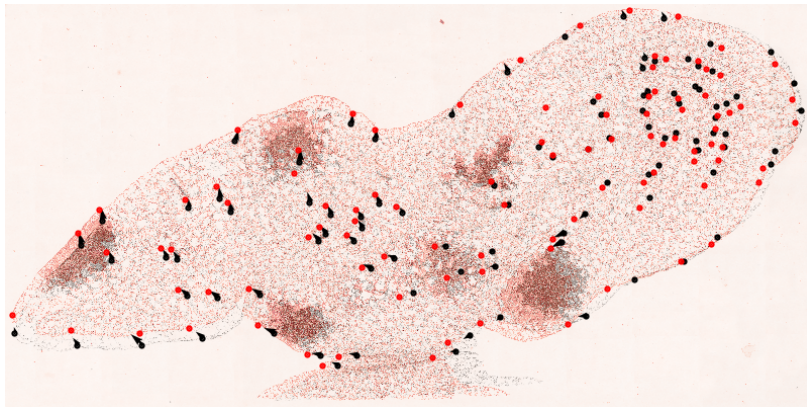


# What makes an application HPC?

Anything that your laptop can't do on its own...

- ▶ More compute power
- ▶ More memory
- ▶ Specialist hardware
- ▶ Large dataset storage
- ▶ Better visualisation
- ▶ etc...

# A real world example — image registration



# How do I write code for an HPC?

Work out what your code needs:

Large Memory - A bigger node than your laptop

- ▶ Code is fast enough, but need more memory
- ▶ No changes to “laptop code”
- ▶ Use HPC large memory nodes for larger datasets

High Throughput - Run many jobs at once

- ▶ Code is fast enough, but need to run lots
- ▶ No changes to “laptop code”
- ▶ Use HPC to run many jobs in parallel

# How do I write code for an HPC?

Work out what your code needs:

## Shared Memory - Multiple cores on one node

- ▶ Code fits in a single node's memory, but needs to be faster
- ▶ Parallel algorithms for “bottlenecks” in the code
- ▶ Can often be added to existing serial code

## Message Passing - Multiple cores on multiple nodes

- ▶ Need multiple nodes to fit problem in memory
- ▶ Parallelise data structures and algorithms
- ▶ Serial code usually needs a complete rewrite

# How do I write code for an HPC?

Work out what your code needs:

## GPU Accelerator

- ▶ Code fits in a single node's memory, but needs to be faster
- ▶ Use 1000s of parallel threads for computation
- ▶ Needs large, parallel problems for efficient use of hardware



# You can use parallel libraries to help

## Low level parallel libraries

- ▶ OpenMP - Shared memory helpers for C/C++/Fortran
- ▶ MPI - MPI standard implemented by many vendors, available for many languages
- ▶ CUDA - Run many parallel threads on Nvidia GPUs
- ▶ OpenACC - OpenMP-like primitives for GPU/Intel MIC offload
- ▶ Intel TBB - threading building blocks for parallel applications

No maths in these libraries. Just tools to help you write parallel code.

# But don't reinvent the wheel

## Lots of higher level libraries available

- ▶ General high performance math - GSL, MKL, NAG
- ▶ Linear Algebra - LaPACK, BLAS, Intel MKL
- ▶ ODE/PDE solvers - PETSc, PVODE, FEniCS
- ▶ Scientific frameworks - PETSc, Trilinos
- ▶ Cluster computing - Dask, Hadoop, Spark
- ▶ Domain specific frameworks - Tensorflow, OpenFOAM, MFEM

## Use a dependency manager

- ▶ Give manager a list of required packages and versions
- ▶ Dependency manager coordinates installation
- ▶ Loaded using modules just like system provided software
- ▶ Install multiple versions of a library side-by-side



## Easier installation and reproducibility

- ▶ Easier installation for users
- ▶ Reproducible environment across platforms
- ▶ Easy testing of dependency versions



## pFIRE - Image Registration Code

- ▶ Computationally intensive linear algebra
- ▶ Must scale to very large 3D images — need lots of memory!
- ▶ MPI is the best choice – look for an MPI linear algebra library

## PETSc

- ▶ High performance, parallel linear algebra library
- ▶ Supports MPI or CUDA for parallelism
- ▶ High performance matrix solvers
- ▶ Lots of helpful routines for matrix and vector manipulation

```
Mat A; // Create a 100000x100000 matrix called A
MatCreate(MPI_COMM_WORLD, &A);
MatSetSizes(A, PETSC_DECIDE, PETSC_DECIDE,
            100000, 100000);

Vec X, Y; // Create compatible sized vectors X and Y
MatCreateVecs(A, &X, &Y);

MatSetValue(A, row, col, value, INSERT_VALUES); // Set a matrix value
VecSetValue(Y, loc, value, INSERT_VALUES); // Set a vector value

// Solve the equation AX = Y

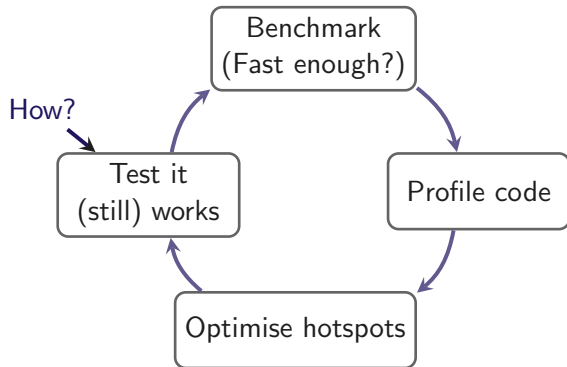
KSP solver; // Create Krylov solver object
KSPCreate(MPI_COMM_WORLD, &solver);
KSPSetOperator(solver, A, A);
KSPSetUp(ksp);
KSPSolve(ksp, Y, X); // Solve for X given M and Y
```

Ignore code performance **until it works!**

- ▶ Test often
- ▶ Focus on hot spots
- ▶ Decide what is “fast enough”

Lots of tools — free and £££

- ▶ Code profilers — Slow sections
- ▶ Memory profilers — Usage and access
- ▶ MPI Profilers — Communication



## Testing early and often helps catch mistakes

- ▶ Test individual functions:
  - ▷ Ensure correct output for valid input
  - ▷ Graceful failures with invalid input
- ▶ Test full program behaviour (integration tests):
  - ▷ Identify useful test cases with known results
  - ▷ Test on different machines/architectures
  - ▷ Regression tests: compare to previous versions

Lots of tools to help automate this

INSIGNEO / pFIRE build passing

Current Branches Build History Pull Requests More options

✓ v0.3.3	add appveyor badges	✓ #80 passed	⌚ 3 min 29 sec	📅 about a month ago
Phil Tooley		↔ c726117		
✓ master	add appveyor badges	✓ #78 passed	⌚ 3 min 59 sec	📅 about a month ago
Phil Tooley		↔ c726117		
✓ master	Merge pr #28 from INSIGNEO	✓ #77 passed	⌚ 3 min 50 sec	📅 about a month ago
Phil Tooley		↔ 73986bd		
✓ develop	enable tests on CI	✓ #75 passed	⌚ 3 min 39 sec	📅 about a month ago
Phil Tooley		↔ efb859c		
✗ develop	fix travis docker cmd	✗ #72 failed	⌚ 4 min 14 sec	📅 about a month ago
Phil Tooley		↔ 94b1fbc		
✗ develop	fix travis docker cmd	✗ #71 failed	⌚ 3 min 36 sec	📅 about a month ago
Phil Tooley		↔ 628ed27		
🕒 develop	fix travis docker cmd	↔ #70 canceled	⌚ 2 min 15 sec	📅 about a month ago
Phil Tooley		↔ bb05517		

## Writing high quality HPC code is not easy, but...

- ▶ Lots of options for applying HPC to a problem
  - determine appropriate approach based on current and future needs
- ▶ Using high-level frameworks provides all the building blocks
  - e.g PETSc for parallel linear algebra, Spark for big data
- ▶ Small changes can bring big gains
  - parallelising just the slowest part can really help with speed
- ▶ Software development best practice helps create robust software
  - regular testing finds bugs early, version control makes fixing them easier

With the right tools, *anyone* can write good HPC code







Thank You

- ▶ Slides: [https://github.com/ptooley/hpc\\_dragons](https://github.com/ptooley/hpc_dragons)
- ▶ RSE Code Clinic:  
Free 30 minute advice sessions: <https://rse.shef.ac.uk/support/code-clinic/>
- ▶ Best Practice Guidance  
RSE@Sheffield: <https://rse.shef.ac.uk>  
SSI: <https://software.ac.uk>
- ▶ HPC Programming  
PRACE Training: <http://www.training.prace-ri.eu/>  
LLNL Online Training: <https://hpc.llnl.gov/training/tutorials>  
GPU@Sheffield: <http://gpucomputing.shef.ac.uk/>  
ShARC Documentation: <http://docs.hpc.shef.ac.uk/en/latest/sharc/>



# COMPBIOMED CONFERENCE 2019

[www.compbio-med-conference.org](http://www.compbio-med-conference.org)

25-27 SEPTEMBER 2019

IET London  
2 Savoy Place. London,  
WC2R 0BL

Conference Chair  
Prof Peter Coveney



## Plenary Speakers

- Andrew Hopkins (Exscientia)
- William L. Jorgensen (Yale)
- Amanda Randles (Durham, NC)
- Anne M. Robertson (Pittsburgh)
- Oliver Röhrle (Stuttgart)

## Conference Themes

### Biomedical Applications (including)

Cardiac & Bloodflow, Neuromusculoskeletal,  
Brain  
Molecular Medicine

### Methodology (including)

Uncertainty Quantification  
Role of Theory, Modelling and Simulation in  
Biomedicine

### Technology & Outreach (including)

Education, Training & Public Awareness  
Regulatory Science and *in silico* Trials